ENOX DIGITAL PUBLISHING - VERIFIED LICENSED EDITION
Read. Learn. Grow.

Title: Production Inference: Low-Latency Language & Multimodal Systems
Author: Nadia Lindqvist & Rohan Deshmukh
ISBN: 978-81-974021-3-8

Chapter 1: The Memory Wall in Autoregressive Decoding

Token generation is overwhelmingly memory-bandwidth bound rather than compute bound. Understanding arithmetic intensity lets you predict latency before provisioning a single GPU cluster.

Licensed for personal reading via ENOX Secure Digital Delivery.